

## ▼ Ishuffle Music Recommendation system

In this project we would be implementing a music recommendation system.

### Import Packages

```
# Core data manipulation packages
import pandas as pd
import numpy as np

# For our EDA charts
import matplotlib.pyplot as plt
import matplotlib.ticker as ticker
import seaborn as sns

# For creating sparse matrix
from scipy.sparse import csr_matrix
from scipy.sparse import coo_matrix

# use K-Nearest Neighbors to find cosine distance amongst songs
from sklearn.neighbors import NearestNeighbors
from sklearn.model_selection import train_test_split

# use decomposition in our matrix factorization
import sklearn
from sklearn.decomposition import TruncatedSVD

# set seed for reproducibility of random number initializations
seed = np.random.RandomState(seed=42)
```

### Load Data

In this section we read data from all our files

```
#load our files
#training files
triplets_file = 'https://static.turi.com/datasets/millionsong/10000.txt'
songs_metadata_file = 'https://static.turi.com/datasets/millionsong/song_data.csv'

#test file
test_df = pd.read_csv('test_data.csv')

#create first dataframe
song_df_1 = pd.read_table(triplets_file,header=None)
song_df_1.columns = ['user_id', 'song_id', 'listen_count']

#Read song file and create second Dataframe
song_df_2 = pd.read_csv(songs_metadata_file)
```

```
#Merge the two dataframes above to create input dataframe for recommender systems
train_df = pd.merge(song_df_1, song_df_2.drop_duplicates(['song_id']), on="song_id")
```

## Feature Engineering

```
#Merge song title and artist_name columns to make a new feature 'song' which is wha
train_df ['song'] = train_df ['title'].map(str) + " - " + train_df['artist_name']
```

## Remove Duplicates from our Dataframe

```
if not train_df[train_df.duplicated(['user_id', 'song'])].empty:
    initial_rows = train_df.shape[0]
    print('Initial dataframe shape {0}'.format(train_df.shape))
    train_df = train_df.drop_duplicates(['user_id', 'song'])
    current_rows = train_df.shape[0]
    print('New dataframe shape {0}'.format(train_df.shape))
    print('Removed {0} rows'.format(initial_rows - current_rows))
```

## Exploratory Data Analysis(EDA)

### Showing the most popular songs in the dataset

```
print("# Top Songs with most total plays:")
#for all user that have played a song and calculate the total sum of the listen co
song_rank = (train_df.groupby(['song']).agg({'user_id': 'count', 'listen_count': 'sum
print(song_rank.head(20))
```

## Visualize Data

We use a barchart to depict the top 20 most listened to songs

```
# our standard bar chart in a function below
def bar_chart_int(x,y,x_label,y_label,title,caption,total_val):
    fig, ax = plt.subplots()
    fig.set_size_inches(16, 5)
    ax = sns.barplot(x[:20], y[:20], palette="PuRd")
    ax.set_xlabel(x_label,fontweight='bold')
    ax.set_ylabel(y_label,fontweight='bold')
    ax.set_title(title,fontweight='bold')
    ax.get_yaxis().set_major_formatter(ticker.FuncFormatter(lambda x, p: '{:,}'.fo

# our bar label placement
for p in ax.patches:
    height = p.get_height()
    pct = 100*(height/total_val)
    ax.text(p.get_x()+p.get_width()/2.,
            height + 3,
            '{:1.1f}%'.format(pct),
            ha="center",verticalalignment='bottom',color='black')
```

```

# our caption statement
ax.text(19, max(y[:20])*0.95, caption, style='italic',fontsize=12,horizontalal:

plt.xticks(rotation=90)
plt.show()

#create a bar chart showing the ranking of top 20 most played songs
c1 = song_rank
x = c1.index
y = c1.userSongPlays
x_label = 'Song Name'
y_label = 'Listen Count'
title = 'Total plays per Song'
caption = 'Percentages are of song plays'
total_val = c1.userSongPlays.sum()
bar_chart_int(x,y,x_label,y_label,title,caption,total_val)

```

Function to get all songs listened to by a user in the dataset

```

def get_user_songs(user):
    user_data = train_df[train_df['user_id'] == user]
    user_items = list(user_data['song'].unique())
    return user_items

#display songs a randomly selected user has listened to in the dataset
index = seed.choice(train_df.shape[0])
uid = train_df.iloc[index]['user_id']
print("User {} has listened to these songs:".format(uid))
for i,song in enumerate(get_user_songs(uid)):
    print("{} . {}".format(i,song))

def get_actual(user):
    user_data = test_df[test_df['user_id'] == user]
    user_items = list(user_data['songs'].unique())
    return user_items

```

## ▼ Model Creation

### Model 1: Popularity based (Baseline model)

Our first model is the Popularity-based model, this model recommends music to users based on th

### Popularity-based recommender model

```

#Use the popularity based recommender system model to make recommendations
def popularity_recommend(user_id):
    popularity_rank = (train_df.groupby(['song']).agg({'user_id':'count','listen
    popularity_rank['Rank'] = popularity_rank['popularity_count'].rank(ascending=
    #Get the top 10 recommendations

```

```

popularity_recommendations = popularity_rank.head(10)

user_recommendations = popularity_recommendations

#Add user_id column for which the recommendations are being generated
user_recommendations['user_id'] = user_id
#Bring user_id column to the front
cols = user_recommendations.columns.tolist()
cols = cols[-1:] + cols[:-1]
user_recommendations = user_recommendations[cols]

return user_recommendations

```

Use the popularity model to make some predictions

```

print("-----")
print("Recommendation based on the most popular songs")
print("-----")
popularity_recommend("4bd88bfb25263a75bbdd467e74018f4ae570e5df")

```

## Model 2: K-Nearest Neighbors

As our first iteration of a basic collaborative recommender, we will build a sparse matrix comparing data will then be passed through a latent mapping algorithm, K-nearest neighbors, to determine correlations. This will help us determine which songs are most similar as in shortest distance apart.

### Prepare Sparse Matrix

In this section, we fit data into a sparse matrix of songs (row) vs. user (column). This matrix captures with number of listens in each respective cell.

```

# function to fit dataframe into a sparse matrix of song name (row) vs user (column)
# in terms of listen count
def data_to_sparse(data, index, columns, values):
    pivot = data.pivot(index=index, columns=columns, values=values).fillna(0)
    sparse = csr_matrix(pivot.values)
    print(sparse.shape)
    return pivot, sparse

```

Use K Nearest Neighbors to determine cosine distance amongst songs

```

def fit_knn(sparse):
    knn = NearestNeighbors(metric='cosine')
    knn.fit(sparse)
    print(knn)
    return knn

```

### Create Sparse Matrix using our DataFrame

```

pivot_df, sparse_df = data.co_sparse(index=[user_id, song], columns=[user_id, song], values=listen_count)

```

```

↳ (9953, 76353)

```

```

pivot_df.head(10)

```

Fit our sparse matrix to our knn model

```

knn = fit_knn(sparse_df)

```

In this function we make our recommendations based on our knn model, we lookup song similarity matrix, with cosine distance in parentheses.

```

def knn_recommend(user_id, data, song, model, k):
    distances, indices = (model.kneighbors(data.loc[song].values.reshape(1,-1), n_neighbors=k))
    predicted = []
    for i in range(0, len(distances.flatten())):
        if i == 0:
            print(("KNN Song Recommendations for user {} cause they like '{}':\n".format(user_id, song)))
        else:
            print('{}: {} ({:.3f})'.format(i, data.index[indices.flatten()[i]], distances.flatten()[i]))
            predicted.append(data.index[indices.flatten()[i]])
    return predicted

```

### Model 3 Matrix Factorization

Similar with kNN, we convert our training data into a 2D matrix (called a utility matrix here) and fill user\_id and our column is song

```

pivot_df2 = train_df.pivot(index = 'user_id', columns = 'song', values = 'listen_count')
pivot_df2.head()

```

We then transpose this utility matrix, so that the songs become rows and user\_id become columns:

```

X = pivot_df2.values.T
X.shape

```

```

↳ (9953, 76353)

```

After using TruncatedSVD to decompose it, we fit it into the model for dimensionality reduction. The number of columns since we must preserve the song names. We choose n\_components = 12 for just 12 latent dimensions have been reduced significantly from 9953 X 76353 to 9953 X 12

```

SVD = TruncatedSVD(n_components=12, random_state=17)
matrix = SVD.fit_transform(X)
matrix.shape

```



We calculate the Pearson's R correlation coefficient for every song pair in our final matrix.

```
correlation = np.corrcoef(matrix)
correlation.shape
```



We recommend songs using this function. Given a song we find all songs that have high correlation with them.

```
def matrix_factorization_recommend(user_id, song):
    songs = pivot_df2.columns
    song_list = list(songs)
    song_index = song_list.index(song)
    correlation_song_index = correlation[song_index]
    print("\nMatrix Factorization recommendation for user {} because they like '{}':".format(user_id, song))
    for i, song in enumerate(list(songs[(correlation_song_index > 0.989)])):
        print("{}: {}".format(i, song))
    return list(songs[(correlation_song_index > 0.989)])
```

## Model Evaluation

Make predictions using our two models, calculate and compare the precision and recall of the two

```
user_id = "4bd88bfb25263a75bbdd467e74018f4ae570e5df"
user_songs = get_user_songs(user_id)
actual_songs = get_actual(user_id)
#knn recommended
knn_predicted = knn_recommend(user_id, pivot_df, user_songs[5], knn, 10)
#Matrix factorization recommended
matrix_fac_predicted = matrix_factorization_recommend(user_id, user_songs[5])
```



We define two functions to calculate the precision and recall for our models.

```
def calc_precision(predicted, actual):
    prec = [value for value in predicted if value in actual]
    prec = np.round(float(len(prec)) / float(len(predicted)), 4) #tp/tp+fp
    return prec

def calc_recall(predicted, actual):
    reca = [value for value in predicted if value in actual]
    reca = np.round(float(len(reca)) / float(len(actual)), 4) #tp/tp+fn
    return reca
```

Calculate Precision and Recall for KNN

```
knn_prec = calc_precision(knn_predicted, actual_songs)
knn_rec = calc_recall(knn_predicted, actual_songs)
```

Calculate Precision and Recall for Matrix Factorization

```
maf_prec = calc_precision(matrix_fac_predicted, actual_songs)
maf_rec = calc_recall(matrix_fac_predicted, actual_songs)
```

Compare Precision and Recall for Both Models

```
if knn_prec > maf_prec:
    print("knn has higher precision of {} compared to the precision of matrix factor:
    print("-----")
else:
    print("MAF has higher precision of {} compared to the precision of KNN which is .
    print("-----")

if knn_rec > maf_rec:
    print("knn has higher recall of {} compared to the recall of matrix factorization
    print("-----")
else:
    print("MAF has higher precision of {} compared to the recall of KNN which is {}".
    print("-----")
```



## ▼ Conclusion

After evaluating our models we can see that knn had a higher precision and recall compared to the

### Demo

```
#randomly select a user from the dataset
index = seed.choice(train_df.shape[0])
uid = train_df.iloc[index]['user_id']
#get songs userhas listened to
user_songs= get_user_songs(uid)
song_index = seed.choice(len(user_songs))
print("User {} has listened to this song '{}'" therefore they will like the following
#knn recommend similar songs
knn_predicted = knn_recommend(uid, pivot_df,user_songs[song_index],knn,10)
#Matrix factorization recommend similar songs
matrix_fac_predicted = matrix_factorization_recommend(uid, user_songs[song_index])
```





